

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: df_iris = pd.read_csv("./iris_csv (1).csv")
```

```
In [4]: df_iris
```

```
Out[4]:
```

	sepalength	sepalwidth	petallength	petalwidth	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...	...	...	...	...	...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 5 columns

```
In [5]: # Drop unnecessart rows
df_iris = df_iris.drop(index = [2, 18])

# Drop rows with na values
df_iris = df_iris.dropna()
```

```
In [6]: df_iris.shape
```

```
Out[6]: (128, 5)
```

```
In [7]: df_iris["class"].value_counts()
```

```
Out[7]: class
Iris-versicolor    50
Iris-virginica     50
Iris-setosa        28
Name: count, dtype: int64
```

## Linear regression

```
In [1]: import statsmodels.api as sm
```

```
In [8]: df_iris.dropna()["sepalwidth"]
```

```
Out[8]: 0      3.5
        1      3.0
        3      3.1
        4      3.6
        5      3.9
        ...
       145     3.0
       146     2.5
       147     3.0
       148     3.4
       149     3.0
        Name: sepalwidth, Length: 128, dtype: float64
```

```
In [9]: # spector_data.exog = sm.add_constant(spector_data.exog, prepend=False)

mod = sm.OLS(endog = df_iris.dropna()["petalwidth"], exog = sm.add_constant(df_iris

res = mod.fit()

print(res.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          petalwidth    R-squared:                0.041
Model:                  OLS          Adj. R-squared:            0.033
Method:                 Least Squares   F-statistic:              5.372
Date:                  Fri, 25 Aug 2023   Prob (F-statistic):       0.0221
Time:                  20:30:54          Log-Likelihood:           -132.08
No. Observations:      128              AIC:                     268.2
Df Residuals:          126              BIC:                     273.9
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025     0.975]
-----
sepalwidth    -0.3588      0.155     -2.318     0.022     -0.665     -0.052
const         2.4370      0.465      5.239     0.000      1.516      3.358
=====
Omnibus:                 14.949    Durbin-Watson:           0.422
Prob(Omnibus):            0.001    Jarque-Bera (JB):         4.846
Skew:                    -0.083    Prob(JB):                 0.0887
Kurtosis:                 2.061    Cond. No.                  25.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [10]: sm.add_constant(df_iris.dropna()["sepalwidth"])
```

Out[10]:

	const	sepalwidth
<b>0</b>	1.0	3.5
<b>1</b>	1.0	3.0
<b>3</b>	1.0	3.1
<b>4</b>	1.0	3.6
<b>5</b>	1.0	3.9
...	...	...
<b>145</b>	1.0	3.0
<b>146</b>	1.0	2.5
<b>147</b>	1.0	3.0
<b>148</b>	1.0	3.4
<b>149</b>	1.0	3.0

128 rows × 2 columns

In [11]: `2.55-(0.41*4.2)`

Out[11]: 0.8279999999999998

In [12]: `sm.add_constant(pd.DataFrame({"c": [1], "a": [4.2]}))`

Out[12]:

	c	a
<b>0</b>	1	4.2

In [13]: `pd.DataFrame({"c": [1], "a": [4.2]})`

Out[13]:

	c	a
<b>0</b>	1	4.2

In [14]: `res.predict(pd.DataFrame({"c": [1], "a": [4.2]}))`

Out[14]: 0 9.876792  
dtype: float64

## kNN

In [ ]: `!pip install scikit-learn`

In [15]: `from sklearn.neighbors import KNeighborsClassifier`

```
In [16]: model_knn = KNeighborsClassifier(n_neighbors = 5)
```

```
In [17]: df_iris.iloc[:, :-1]
```

```
Out[17]:
```

	sepalength	sepalwidth	petallength	petalwidth
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
5	5.4	3.9	1.7	0.4
...	...	...	...	...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

128 rows × 4 columns

```
In [18]: # Fitting / training the kNN machine learning classifier

model_knn.fit(X = df_iris.iloc[:, :-1], y = df_iris["class"])
```

```
Out[18]:
```

▼ KNeighborsClassifier

KNeighborsClassifier()

```
In [19]: # Predicting the class label for a new datapoint

model_knn.predict([[1.2, 1, 3, 4.9]])
```

```
C:\python 3114\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names
warnings.warn(
```

```
Out[19]: array(['Iris-versicolor'], dtype=object)
```

```
In [ ]:
```

```
In [ ]:
```

## Train -test split

```
In [20]: df_iris
```

Out[20]:

	sepalength	sepalwidth	petallength	petalwidth	class
<b>0</b>	5.1	3.5	1.4	0.2	Iris-setosa
<b>1</b>	4.9	3.0	1.4	0.2	Iris-setosa
<b>3</b>	4.6	3.1	1.5	0.2	Iris-setosa
<b>4</b>	5.0	3.6	1.4	0.2	Iris-setosa
<b>5</b>	5.4	3.9	1.7	0.4	Iris-setosa
<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>
<b>145</b>	6.7	3.0	5.2	2.3	Iris-virginica
<b>146</b>	6.3	2.5	5.0	1.9	Iris-virginica
<b>147</b>	6.5	3.0	5.2	2.0	Iris-virginica
<b>148</b>	6.2	3.4	5.4	2.3	Iris-virginica
<b>149</b>	5.9	3.0	5.1	1.8	Iris-virginica

128 rows × 5 columns

## Regression

```
In [21]: # Target (dependent variable) : petalwidth  
# Independednt variables : sepalength, sepalwidth, petallength
```

### Linear train test split

```
In [22]: df_train = df_iris.iloc[:int(128*0.8), :]  
df_test = df_iris.iloc[int(128*0.8):, :]
```

```
In [23]: print(df_train.shape)  
print(df_test.shape)
```

(102, 5)

(26, 5)

### Random train test split

```
In [24]: np.random.choice(range(128), int(128*0.8), replace = False)
```

```
Out[24]: array([ 34, 110, 55, 7, 57, 112, 78, 65, 40, 87, 106, 95, 71,
                58, 122, 52, 88, 121, 93, 25, 36, 105, 2, 6, 39, 90,
                99, 114, 43, 83, 84, 92, 10, 32, 15, 33, 22, 79, 8,
                81, 11, 61, 16, 80, 103, 62, 54, 66, 45, 41, 98, 127,
                82, 38, 59, 75, 48, 49, 89, 96, 85, 100, 125, 9, 86,
                3, 63, 47, 50, 13, 51, 53, 14, 28, 94, 19, 108, 69,
                64, 104, 26, 12, 109, 118, 0, 29, 44, 4, 31, 68, 72,
                42, 73, 117, 107, 119, 23, 101, 5, 60, 97, 67])
```

```
In [25]: # Generate random indices for the training data
df_train = df_iris.iloc[np.random.choice(range(128), int(128*0.8), replace = False)]
df_train
```

```
Out[25]:
```

	sepalength	sepalwidth	petallength	petalwidth	class
<b>93</b>	5.0	2.3	3.3	1.0	Iris-versicolor
<b>68</b>	6.2	2.2	4.5	1.5	Iris-versicolor
<b>115</b>	6.4	3.2	5.3	2.3	Iris-virginica
<b>21</b>	5.1	3.7	1.5	0.4	Iris-setosa
<b>57</b>	4.9	2.4	3.3	1.0	Iris-versicolor
...	...	...	...	...	...
<b>118</b>	7.7	2.6	6.9	2.3	Iris-virginica
<b>130</b>	7.4	2.8	6.1	1.9	Iris-virginica
<b>49</b>	5.0	3.3	1.4	0.2	Iris-setosa
<b>83</b>	6.0	2.7	5.1	1.6	Iris-versicolor
<b>91</b>	6.1	3.0	4.6	1.4	Iris-versicolor

102 rows × 5 columns

```
In [26]: # Get those indices which has not been used in training
set(df_iris.index) - set(df_train.index)
```

```
Out[26]: {19,  
          25,  
          39,  
          42,  
          46,  
          50,  
          64,  
          82,  
          85,  
          86,  
          87,  
          88,  
          89,  
          92,  
          101,  
          103,  
          111,  
          112,  
          121,  
          127,  
          129,  
          138,  
          140,  
          147,  
          148,  
          149}
```

```
In [27]: # Use the left out indices for test data  
df_test = df_iris.loc[list(set(df_iris.index) - set(df_train.index)), :]  
df_test
```

Out[27]:

	sepalength	sepalwidth	petallength	petalwidth	class
129	7.2	3.0	5.8	1.6	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
140	6.7	3.1	5.6	2.4	Iris-virginica
19	5.1	3.8	1.5	0.3	Iris-setosa
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica
25	5.0	3.0	1.6	0.2	Iris-setosa
39	5.1	3.4	1.5	0.2	Iris-setosa
42	4.4	3.2	1.3	0.2	Iris-setosa
46	5.1	3.8	1.6	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor
101	5.8	2.7	5.1	1.9	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica

## Using train\_test\_split function from sklearn

```
In [28]: X = df_iris.loc[:, ["sepalength", "sepalwidth", "petallength"]]  
         y = df_iris.loc[:, "petalwidth"]
```

```
In [29]: X
```



Out[29]:

	sepalength	sepalwidth	petallength
0	5.1	3.5	1.4
1	4.9	3.0	1.4
3	4.6	3.1	1.5
4	5.0	3.6	1.4
5	5.4	3.9	1.7
...	...	...	...
145	6.7	3.0	5.2
146	6.3	2.5	5.0
147	6.5	3.0	5.2
148	6.2	3.4	5.4
149	5.9	3.0	5.1

128 rows × 3 columns

In [30]: `y`

Out[30]:

0	0.2
1	0.2
3	0.2
4	0.2
5	0.4
...	...
145	2.3
146	1.9
147	2.0
148	2.3
149	1.8

Name: petalwidth, Length: 128, dtype: float64

In [31]: `from sklearn.model_selection import train_test_split`  
`X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.8, random_`

In [32]: `X_train`

```
Out[32]:
```

	sepalength	sepalwidth	petallength
92	5.8	2.6	4.0
100	6.3	3.3	6.0
69	5.6	2.5	3.9
0	5.1	3.5	1.4
22	4.6	3.6	1.0
...	...	...	...
128	6.4	2.8	5.6
24	4.8	3.4	1.9
114	5.8	2.8	5.1
73	6.1	2.8	4.7
124	6.7	3.3	5.7

102 rows × 3 columns

```
In [33]: X_test.shape
```

```
Out[33]: (26, 3)
```

```
In [34]: y_train
```

```
Out[34]: 92      1.2
100     2.5
69      1.1
0       0.2
22      0.2
...
128     2.1
24      0.2
114     2.4
73      1.2
124     2.1
Name: petalwidth, Length: 102, dtype: float64
```

```
In [35]: len(y_test)
```

```
Out[35]: 26
```

## Model training

```
In [36]: from sklearn.linear_model import LinearRegression
```

```
In [37]: # Creating the model object
model_lr = LinearRegression()
```

```
In [38]: model_lr.fit(X_train, y_train)
```

```
Out[38]: ▼ LinearRegression  
LinearRegression()
```

```
In [39]: model_lr.score(X_train, y_train)
```

```
Out[39]: 0.9104513929181576
```

```
In [40]: dir(model_lr)
```

```
Out[40]: ['__abstractmethods__',
          '__annotations__',
          '__class__',
          '__delattr__',
          '__dict__',
          '__dir__',
          '__doc__',
          '__eq__',
          '__format__',
          '__ge__',
          '__getattr__',
          '__getstate__',
          '__gt__',
          '__hash__',
          '__init__',
          '__init_subclass__',
          '__le__',
          '__lt__',
          '__module__',
          '__ne__',
          '__new__',
          '__reduce__',
          '__reduce_ex__',
          '__repr__',
          '__setattr__',
          '__setstate__',
          '__sizeof__',
          '__sklearn_clone__',
          '__str__',
          '__subclasshook__',
          '__weakref__',
          '_abc_impl',
          '_build_request_for_signature',
          '_check_feature_names',
          '_check_n_features',
          '_decision_function',
          '_estimator_type',
          '_get_default_requests',
          '_get_metadata_request',
          '_get_param_names',
          '_get_tags',
          '_more_tags',
          '_parameter_constraints',
          '_repr_html_',
          '_repr_html_inner',
          '_repr_mimebundle_',
          '_set_intercept',
          '_validate_data',
          '_validate_params',
          'coef_',
          'copy_X',
          'feature_names_in_',
          'fit',
          'fit_intercept',
          'get_metadata_routing',
          'get_params',
```

```
'intercept_',  
'n_features_in_',  
'n_jobs',  
'positive',  
'predict',  
'rank_',  
'score',  
'set_fit_request',  
'set_params',  
'set_score_request',  
'singular_']
```

```
In [41]: model_lr.feature_names_in_
```

```
Out[41]: array(['sepallength', 'sepalwidth', 'petallength'], dtype=object)
```

```
In [42]: model_lr.coef_
```

```
Out[42]: array([-0.26107413,  0.24827545,  0.54700606])
```

```
In [43]: model_lr.intercept_
```

```
Out[43]: -0.09575419984582423
```

```
In [44]: dict(zip(model_lr.feature_names_in_, model_lr.coef_))
```

```
Out[44]: {'sepallength': -0.2610741299844725,  
          'sepalwidth': 0.24827545186904953,  
          'petallength': 0.5470060612404106}
```

```
In [45]: # Prediction on test data  
y_pred_lr1 = model_lr.predict(X_test)  
y_pred_lr1
```

```
Out[45]: array([1.63490579, 1.07203126, 0.01155363, 1.22739587, 2.0152503 ,  
                1.51932711, 1.57274784, 1.60638655, 2.04271153, 0.25968118,  
                1.96062365, 0.18399263, 1.49563152, 0.39263006, 2.313833 ,  
                1.97553833, 1.41738323, 2.02278159, 1.79652183, 1.70719847,  
                0.17894711, 0.26849234, 1.98793697, 0.311896 , 1.73684966,  
                1.30308442])
```

```
In [46]: y_test.values
```

```
Out[46]: array([1.7, 1. , 0.3, 1.3, 2.3, 1.5, 1.4, 1.8, 1.8, 0.2, 1.8, 0.2, 1.5,  
                0.4, 2.3, 2.3, 1.3, 2.1, 2. , 1.7, 0.3, 0.2, 2.4, 0.4, 1.8, 1. ])
```

```
In [47]: y_test.shape
```

```
Out[47]: (26,)
```

```
In [49]: np.sqrt(np.sum((y_test.values - y_pred_lr1)**2)/26) # RMSE
```

```
Out[49]: 0.17482930839177857
```

```
In [50]: np.sum((y_test.values - y_pred_lr1)**2)/26 #MSE
```

Out[50]: 0.030565287072747614

```
In [51]: from sklearn.metrics import mean_squared_error  
mean_squared_error(y_test.values, y_pred_lr1)
```

Out[51]: 0.030565287072747614

```
In [52]: np.round(y_pred_lr1, 1)
```

Out[52]: array([1.6, 1.1, 0. , 1.2, 2. , 1.5, 1.6, 1.6, 2. , 0.3, 2. , 0.2, 1.5,  
0.4, 2.3, 2. , 1.4, 2. , 1.8, 1.7, 0.2, 0.3, 2. , 0.3, 1.7, 1.3])

```
In [53]: TSS = np.sum((y_test.values - np.round(y_test.values.mean(), 1))**2)  
ESS = np.sum((np.round(y_pred_lr1, 1) - np.round(y_test.values.mean(), 1))**2)  
RSS = np.sum((y_test.values - np.round(y_pred_lr1, 1))**2)
```

```
In [54]: TSS
```

Out[54]: 14.18

```
In [55]: ESS + RSS
```

Out[55]: 13.66

```
In [ ]:
```

## Classification using kNN

- Feature set (independent variables): 'sepalength', 'sepalwidth', 'petallength', 'petalwidth'
- Target class (dependent variable): 'class'

```
In [56]: X = df_iris.loc[:, ['sepalength', 'sepalwidth', 'petallength', 'petalwidth']]  
y = df_iris.loc[:, 'class']
```

```
In [57]: X
```

```
Out[57]:
```

	sepalength	sepalwidth	petallength	petalwidth
<b>0</b>	5.1	3.5	1.4	0.2
<b>1</b>	4.9	3.0	1.4	0.2
<b>3</b>	4.6	3.1	1.5	0.2
<b>4</b>	5.0	3.6	1.4	0.2
<b>5</b>	5.4	3.9	1.7	0.4
<b>...</b>	...	...	...	...
<b>145</b>	6.7	3.0	5.2	2.3
<b>146</b>	6.3	2.5	5.0	1.9
<b>147</b>	6.5	3.0	5.2	2.0
<b>148</b>	6.2	3.4	5.4	2.3
<b>149</b>	5.9	3.0	5.1	1.8

128 rows × 4 columns

```
In [58]: y
```

```
Out[58]: 0      Iris-setosa
1      Iris-setosa
3      Iris-setosa
4      Iris-setosa
5      Iris-setosa
...
145    Iris-virginica
146    Iris-virginica
147    Iris-virginica
148    Iris-virginica
149    Iris-virginica
Name: class, Length: 128, dtype: object
```

## Performing train test split

```
In [98]: from sklearn.model_selection import train_test_split

X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, train_size = 0.8, ran
```

```
In [99]: X_train2
```

Out[99]:

	sepallength	sepalwidth	petallength	petalwidth
143	6.8	3.2	5.9	2.3
53	5.5	2.3	4.0	1.3
89	5.5	2.5	4.0	1.3
88	5.6	3.0	4.1	1.3
124	6.7	3.3	5.7	2.1
...	...	...	...	...
78	6.0	2.9	4.5	1.5
76	6.8	2.8	4.8	1.4
139	6.9	3.1	5.4	2.1
114	5.8	2.8	5.1	2.4
148	6.2	3.4	5.4	2.3

102 rows × 4 columns

In [100...

X\_test2



Out[100]:

	sepalength	sepalwidth	petallength	petalwidth
<b>17</b>	5.1	3.5	1.4	0.3
<b>4</b>	5.0	3.6	1.4	0.2
<b>61</b>	5.9	3.0	4.2	1.5
<b>107</b>	7.3	2.9	6.3	1.8
<b>20</b>	5.4	3.4	1.7	0.2
<b>123</b>	6.3	2.7	4.9	1.8
<b>43</b>	5.0	3.5	1.6	0.6
<b>65</b>	6.7	3.1	4.4	1.4
<b>93</b>	5.0	2.3	3.3	1.0
<b>106</b>	4.9	2.5	4.5	1.7
<b>137</b>	6.4	3.1	5.5	1.8
<b>116</b>	6.5	3.0	5.5	1.8
<b>47</b>	4.6	3.2	1.4	0.2
<b>57</b>	4.9	2.4	3.3	1.0
<b>42</b>	4.4	3.2	1.3	0.2
<b>118</b>	7.7	2.6	6.9	2.3
<b>132</b>	6.4	2.8	5.6	2.2
<b>97</b>	6.2	2.9	4.3	1.3
<b>50</b>	7.0	3.2	4.7	1.4
<b>90</b>	5.5	2.6	4.4	1.2
<b>92</b>	5.8	2.6	4.0	1.2
<b>48</b>	5.3	3.7	1.5	0.2
<b>104</b>	6.5	3.0	5.8	2.2
<b>45</b>	4.8	3.0	1.4	0.3
<b>125</b>	7.2	3.2	6.0	1.8
<b>7</b>	5.0	3.4	1.5	0.2

In [115... y\_train2

```
Out[115]: 143    Iris-virginica
          53    Iris-versicolor
          89    Iris-versicolor
          88    Iris-versicolor
          124   Iris-virginica
          ...
          78    Iris-versicolor
          76    Iris-versicolor
          139   Iris-virginica
          114   Iris-virginica
          148   Iris-virginica
          Name: class, Length: 102, dtype: object
```

```
In [114... y_test2
```

```
Out[114]: 17    Iris-setosa
          4    Iris-setosa
          61   Iris-versicolor
          107  Iris-virginica
          20   Iris-setosa
          123  Iris-virginica
          43   Iris-setosa
          65   Iris-versicolor
          93   Iris-versicolor
          106  Iris-virginica
          137  Iris-virginica
          116  Iris-virginica
          47   Iris-setosa
          57   Iris-versicolor
          42   Iris-setosa
          118  Iris-virginica
          132  Iris-virginica
          97   Iris-versicolor
          50   Iris-versicolor
          90   Iris-versicolor
          92   Iris-versicolor
          48   Iris-setosa
          104  Iris-virginica
          45   Iris-setosa
          125  Iris-virginica
          7    Iris-setosa
          Name: class, dtype: object
```

```
In [101... y_test2.shape
```

```
Out[101]: (26,)
```

## Train a kNN classifier using the training data

```
In [102... from sklearn.neighbors import KNeighborsClassifier
```

```
In [103... # Initializing the model object
          model_knn = KNeighborsClassifier(n_neighbors = 4)
```

```
In [104... # Fit the model on training data
model_knn.fit(X_train2, y_train2)
```

```
Out[104]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=4)
```

```
In [123... # Run predictions on test data
y_pred_knn = model_knn.predict(X_test2)
y_pred_knn
```

```
Out[123]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',
                'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                'Iris-setosa'], dtype=object)
```

```
In [124... # Original targets
y_test2.values
```

```
Out[124]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-virginica', 'Iris-virginica',
                'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                'Iris-setosa'], dtype=object)
```

```
In [127... incorrects = np.nonzero(model_knn.predict(X_test2) != y_test2)[0]
incorrects
```

```
Out[127]: array([9], dtype=int64)
```

```
In [135... incorrect_indices = np.nonzero(model_knn.predict(X_test2) != y_test2)[0]

print("Length of y_test2:", len(y_test2))
print("Length of incorrect_indices:", len(incorrect_indices))
print("Indices in incorrect_indices:", incorrect_indices)

for index in incorrect_indices:
    if index < len(y_test2):
        original_value = y_test2.iloc[index] # Using iloc to access by integer ind
        predicted_value = model_knn.predict(X_test2.iloc[index].values.reshape(1, -
        print(f"Original Value: {original_value}, Predicted Value: {predicted_value}
    else:
        print(f"Index {index} is out of bounds for y_test2")
```

Length of y\_test2: 26

Length of incorrect\_indices: 1

Indices in incorrect\_indices: [9]

Original Value: Iris-virginica, Predicted Value: Iris-versicolor

```
C:\python 3114\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names
warnings.warn(
```

## Evaluating the kNN classifier

```
In [136... model_knn.score(X_test2, y_test2.values)
```

```
Out[136]: 0.9615384615384616
```

```
In [147... # Confusion matrix for our test results
from sklearn.metrics import confusion_matrix

confusion_matrix(y_test2.values, y_pred_knn)
```

```
Out[147]: array([[9, 0, 0],
                [0, 8, 0],
                [0, 1, 8]], dtype=int64)
```

```
In [138... 25/26
```

```
Out[138]: 0.9615384615384616
```

```
In [139... from sklearn.metrics import accuracy_score
accuracy_score(y_test2, y_pred_knn)
```

```
Out[139]: 0.9615384615384616
```

```
In [144... from sklearn.metrics import precision_score
precision_score(y_test2, y_pred_knn, average="weighted")
```

```
Out[144]: 0.9658119658119658
```

## Classification using Logistic regression

- Feature set (independent variables): 'sepalength', 'sepalwidth', 'petallength', 'petalwidth'
- Target class (dependent variable): 'class'

```
In [113... from sklearn.linear_model import LogisticRegression
```

```
In [176... from sklearn.model_selection import train_test_split
```

```
X_train3, X_test3, y_train3, y_test3 = train_test_split(X, y, train_size = 0.8, ran
```

```
In [303... # Initializing the model object
model_logistic1 = LogisticRegression(random_state = 12).fit(X_train3,y_train3)
```

```
In [304... # Fit the model on training data
model_logistic1.fit(X_train3, y_train3)
```

```
Out[304]: LogisticRegression
LogisticRegression(random_state=12)
```

```
In [305]: # Run predictions on test data
y_pred_lr2 = model_logistic1.predict(X_test3)
y_pred_lr2
```

```
Out[305]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',
                'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                'Iris-setosa'], dtype=object)
```

```
In [306]: # Get probability values on test data
model_logistic1.predict_proba(X_test3)
```

```
Out[306]: array([[9.55403239e-01, 4.45966892e-02, 7.14782839e-08],
                [9.63467651e-01, 3.65323046e-02, 4.47828657e-08],
                [1.89697290e-02, 8.98737115e-01, 8.22931564e-02],
                [1.48171502e-06, 2.18108016e-02, 9.78187717e-01],
                [8.86857701e-01, 1.13142008e-01, 2.90880537e-07],
                [1.06999642e-03, 4.20181301e-01, 5.78748703e-01],
                [9.40671448e-01, 5.93281564e-02, 3.95804966e-07],
                [6.20026320e-03, 9.12660054e-01, 8.11396824e-02],
                [1.40224320e-01, 8.57162866e-01, 2.61281388e-03],
                [1.24958187e-02, 6.20995522e-01, 3.66508660e-01],
                [1.41159992e-04, 1.35692875e-01, 8.64165965e-01],
                [1.16110947e-04, 1.36947749e-01, 8.62936140e-01],
                [9.64239902e-01, 3.57600427e-02, 5.56436090e-08],
                [1.66740364e-01, 8.30700426e-01, 2.55920995e-03],
                [9.75934964e-01, 2.40650057e-02, 3.04402572e-08],
                [1.11054814e-08, 1.11892694e-03, 9.98881062e-01],
                [2.18591285e-05, 3.49853332e-02, 9.64992808e-01],
                [1.07072974e-02, 9.29942834e-01, 5.93498689e-02],
                [2.55652293e-03, 8.45693000e-01, 1.51750477e-01],
                [1.33055040e-02, 9.07123542e-01, 7.95709542e-02],
                [2.25148990e-02, 9.52687865e-01, 2.47972359e-02],
                [9.45765150e-01, 5.42347731e-02, 7.70498752e-08],
                [1.02277379e-05, 2.21428882e-02, 9.77846884e-01],
                [9.46953730e-01, 5.30461636e-02, 1.06780408e-07],
                [9.31431761e-06, 5.05565925e-02, 9.49434093e-01],
                [9.46611609e-01, 5.33883003e-02, 9.07041265e-08]])
```

## Evaluating the Logistic regression classifier

```
In [307]: model_logistic1.score(X_test3, y_test3.values)
```

```
Out[307]: 0.9615384615384616
```

```
In [308]: confusion_matrix(y_test3.values, y_pred_lr2)
```

```
Out[308]: array([[9, 0, 0],
                 [0, 8, 0],
                 [0, 1, 8]], dtype=int64)
```

```
In [312... # Initializing the model object
model_logistic2 = LogisticRegression(random_state = 12,solver='liblinear').fit(X_tr
```

```
In [313... # Fit the model on training data
model_logistic2.fit(X_train3, y_train3)
```

```
Out[313]: ▼ LogisticRegression
LogisticRegression(random_state=12, solver='liblinear')
```

```
In [314... # Run predictions on test data
y_pred_lr3 = model_logistic2.predict(X_test3)
y_pred_lr3
```

```
Out[314]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-virginica', 'Iris-virginica',
                 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                 'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                 'Iris-setosa'], dtype=object)
```

```
In [315... # Get probability values on test data
model_logistic2.predict_proba(X_test3)
```

```
Out[315]: array([[7.91157011e-01, 2.08812269e-01, 3.07195466e-05],
 [7.93944127e-01, 2.06031486e-01, 2.43867568e-05],
 [3.97568242e-02, 6.37740348e-01, 3.22502828e-01],
 [2.28191327e-04, 4.29320145e-01, 5.70451664e-01],
 [7.16118825e-01, 2.83850034e-01, 3.11407430e-05],
 [2.96986768e-03, 4.10284516e-01, 5.86745616e-01],
 [8.34580169e-01, 1.65290863e-01, 1.28968121e-04],
 [4.25757068e-02, 8.47805955e-01, 1.09618338e-01],
 [8.77254436e-02, 7.91997733e-01, 1.20276824e-01],
 [3.15090117e-03, 2.82535775e-01, 7.14313324e-01],
 [1.41443573e-03, 3.42318459e-01, 6.56267105e-01],
 [1.24767063e-03, 3.65982761e-01, 6.32769568e-01],
 [7.48467160e-01, 2.51455239e-01, 7.76008103e-05],
 [1.00964864e-01, 7.70051003e-01, 1.28984133e-01],
 [7.67072192e-01, 2.32841144e-01, 8.66637757e-05],
 [3.11758411e-05, 3.99787882e-01, 6.00180942e-01],
 [4.95300431e-04, 2.70045896e-01, 7.29458803e-01],
 [3.31458945e-02, 7.99774831e-01, 1.67079274e-01],
 [2.76950607e-02, 8.63551784e-01, 1.08753155e-01],
 [1.11143414e-02, 6.14840215e-01, 3.74045444e-01],
 [3.55934778e-02, 7.95358492e-01, 1.69048030e-01],
 [7.87380926e-01, 2.12602719e-01, 1.63540285e-05],
 [4.49519684e-04, 2.55510690e-01, 7.44039790e-01],
 [7.21238090e-01, 2.78668007e-01, 9.39032434e-05],
 [7.28342350e-04, 4.20570451e-01, 5.78701207e-01],
 [7.53648272e-01, 2.46312826e-01, 3.89023051e-05]])
```

```
In [316]: model_logistic2.score(X_test3, y_test3.values)
```

```
Out[316]: 1.0
```

```
In [317]: confusion_matrix(y_test3.values, y_pred_lr3)
```

```
Out[317]: array([[9, 0, 0],
 [0, 8, 0],
 [0, 0, 9]], dtype=int64)
```

## Classification using Support Vector Machine

- Feature set (independent variables): 'sepalwidth', 'sepalwidth', 'petallength', 'petalwidth'
- Target class (dependent variable): 'class'

```
In [182]: from sklearn.svm import SVC
```

```
In [183]: # Initializing the model object
model_svm = SVC()
```

```
In [184]: # Fit the model on training data
model_svm.fit(X_train2, y_train2)
```

Out[184]: 

```
In [185... # Run predictions on test data
y_pred_svc = model_svm.predict(X_test2)
y_pred_svc
```

```
Out[185]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',
                'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                'Iris-setosa'], dtype=object)
```

## Evaluating the SVM classifier

```
In [186... model_svm.score(X_test2, y_test2.values)
```

Out[186]: 0.9615384615384616

```
In [187... # Confusion matrix for our test results
from sklearn.metrics import confusion_matrix

confusion_matrix(y_test2.values, y_pred_svc)
```

```
Out[187]: array([[9, 0, 0],
                [0, 8, 0],
                [0, 1, 8]], dtype=int64)
```

```
In [193... # Using linear kernel instead of rbf

# Initializing the model object
model_svm2 = SVC(kernel = "linear")

# Fit the model on training data
model_svm2.fit(X_train2, y_train2)

# Run predictions on test data
y_pred_svc2 = model_svm2.predict(X_test2)
y_pred_svc2
```

```
Out[193]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-virginica', 'Iris-virginica',
                'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                'Iris-setosa'], dtype=object)
```

```
In [194... model_svm2.score(X_test2, y_test2.values)
```



```
Out[194]: 1.0
```

```
In [195... confusion_matrix(y_test2.values, y_pred_svc2)
```

```
Out[195]: array([[9, 0, 0],
                 [0, 8, 0],
                 [0, 0, 9]], dtype=int64)
```

## Classification using Decision Tree classifier

- Feature set (independent variables): 'sepalwidth', 'sepalwidth', 'petallength', 'petalwidth'
- Target class (dependent variable): 'class'

```
In [196... from sklearn.tree import DecisionTreeClassifier
```

```
In [246... # Initializing the model object
model_dt = DecisionTreeClassifier(random_state = 10)
```

```
In [247... # Fit the model on training data
model_dt.fit(X_train2, y_train2)
```

```
Out[247]: ▼      DecisionTreeClassifier
          DecisionTreeClassifier(random_state=10)
```

```
In [248... # Run predictions on test data
y_pred_dt = model_dt.predict(X_test2)
y_pred_dt
```

```
Out[248]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',
                 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                 'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                 'Iris-setosa'], dtype=object)
```

## Evaluating the Decision Tree classifier

```
In [249... model_dt.score(X_test2, y_test2.values)
```

```
Out[249]: 0.9615384615384616
```

```
In [250... # Confusion matrix for our test results
from sklearn.metrics import confusion_matrix

confusion_matrix(y_test2.values, y_pred_dt)
```

```
Out[250]: array([[9, 0, 0],
                 [0, 8, 0],
                 [0, 1, 8]], dtype=int64)
```

## Classification using Random Forest classifier

- Feature set (independent variables): 'sepalwidth', 'petallength', 'petalwidth'
- Target class (dependent variable): 'class'

```
In [202... from sklearn.ensemble import RandomForestClassifier
```

```
In [203... # Initializing the model object
model_rf = RandomForestClassifier(random_state = 10)
```

```
In [205... # Fit the model on training data
model_rf.fit(X_train2, y_train2)
```

```
Out[205]: ▼      RandomForestClassifier
          RandomForestClassifier(random_state=10)
```

```
In [206... # Run predictions on test data
y_pred_rf = model_rf.predict(X_test2)
y_pred_rf
```

```
Out[206]: array(['Iris-setosa', 'Iris-setosa', 'Iris-versicolor', 'Iris-virginica',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-versicolor', 'Iris-virginica',
                 'Iris-virginica', 'Iris-setosa', 'Iris-versicolor', 'Iris-setosa',
                 'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
                 'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
                 'Iris-setosa', 'Iris-virginica', 'Iris-setosa', 'Iris-virginica',
                 'Iris-setosa'], dtype=object)
```

## Evaluating the Random Forest classifier

```
In [207... model_rf.score(X_test2, y_test2.values)
```

```
Out[207]: 0.9615384615384616
```

```
In [208... model_rf.get_params()
```

```
Out[208]: {'bootstrap': True,
           'ccp_alpha': 0.0,
           'class_weight': None,
           'criterion': 'gini',
           'max_depth': None,
           'max_features': 'sqrt',
           'max_leaf_nodes': None,
           'max_samples': None,
           'min_impurity_decrease': 0.0,
           'min_samples_leaf': 1,
           'min_samples_split': 2,
           'min_weight_fraction_leaf': 0.0,
           'n_estimators': 100,
           'n_jobs': None,
           'oob_score': False,
           'random_state': 10,
           'verbose': 0,
           'warm_start': False}
```

```
In [209... confusion_matrix(y_test2.values, y_pred_rf)
```

```
Out[209]: array([[9, 0, 0],
                 [0, 8, 0],
                 [0, 1, 8]], dtype=int64)
```

# Unsupervised ML

## k-means

```
In [210... from sklearn.cluster import KMeans
```

```
In [287... # Initializing the model object
model_kmeans = KMeans(n_clusters = 3, random_state=0, n_init='auto')
```

```
In [288... X_train2
```

```
Out[288]:
```

	sepalength	sepalwidth	petallength	petalwidth
143	6.8	3.2	5.9	2.3
53	5.5	2.3	4.0	1.3
89	5.5	2.5	4.0	1.3
88	5.6	3.0	4.1	1.3
124	6.7	3.3	5.7	2.1
...	...	...	...	...
78	6.0	2.9	4.5	1.5
76	6.8	2.8	4.8	1.4
139	6.9	3.1	5.4	2.1
114	5.8	2.8	5.1	2.4
148	6.2	3.4	5.4	2.3

102 rows × 4 columns

```
In [296... # Fit the model on training data
model_kmeans.fit(X_train2)
```

```
Out[296]: KMeans
KMeans(n_clusters=3, n_init='auto', random_state=0)
```

```
In [297... model_kmeans.labels_
```

```
Out[297]: array([2, 0, 0, 0, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 2, 2,
                0, 1, 0, 0, 2, 0, 1, 2, 2, 1, 2, 0, 0, 1, 0, 0, 0, 0, 0, 0, 2, 2,
                1, 1, 0, 1, 0, 2, 0, 2, 2, 1, 2, 0, 0, 2, 0, 0, 2, 0, 2, 1, 0, 0,
                2, 2, 1, 1, 0, 1, 2, 1, 2, 2, 1, 1, 0, 0, 2, 0, 2, 2, 0, 1, 0, 2,
                0, 0, 2, 0, 0, 0, 2, 0, 1, 0, 0, 2, 0, 2])
```

```
In [298... model_kmeans.predict(X_test2)
```

```
Out[298]: array([1, 1, 0, 2, 1, 0, 1, 0, 0, 0, 2, 2, 1, 0, 1, 2, 2, 0, 0, 0, 0, 1,
                2, 1, 2, 1])
```

## Reverse engineering to validate the clustering

```
In [299... # Creating a 0, 1, 2 mapping dictionary of the actual class labels
y_test_dict = {0: "Iris-versicolor", 1: "Iris-setosa", 2: "Iris-virginica"}
y_test_dict
```

```
Out[299]: {0: 'Iris-versicolor', 1: 'Iris-setosa', 2: 'Iris-virginica'}
```

```
In [300... # List of predicted cluster labels mapped to the dictionary
[y_test_dict[i] for i in model_kmeans.predict(X_test2)]
```

```
Out[300]: ['Iris-setosa',  
           'Iris-setosa',  
           'Iris-versicolor',  
           'Iris-virginica',  
           'Iris-setosa',  
           'Iris-versicolor',  
           'Iris-setosa',  
           'Iris-versicolor',  
           'Iris-versicolor',  
           'Iris-versicolor',  
           'Iris-virginica',  
           'Iris-virginica',  
           'Iris-setosa',  
           'Iris-versicolor',  
           'Iris-setosa',  
           'Iris-virginica',  
           'Iris-virginica',  
           'Iris-versicolor',  
           'Iris-versicolor',  
           'Iris-versicolor',  
           'Iris-versicolor',  
           'Iris-setosa',  
           'Iris-virginica',  
           'Iris-setosa',  
           'Iris-virginica',  
           'Iris-setosa']
```

```
In [301... # Actual test labels  
y_test2
```

```
Out[301]: 17      Iris-setosa
         4      Iris-setosa
         61     Iris-versicolor
         107    Iris-virginica
         20     Iris-setosa
         123    Iris-virginica
         43     Iris-setosa
         65     Iris-versicolor
         93     Iris-versicolor
         106    Iris-virginica
         137    Iris-virginica
         116    Iris-virginica
         47     Iris-setosa
         57     Iris-versicolor
         42     Iris-setosa
         118    Iris-virginica
         132    Iris-virginica
         97     Iris-versicolor
         50     Iris-versicolor
         90     Iris-versicolor
         92     Iris-versicolor
         48     Iris-setosa
         104    Iris-virginica
         45     Iris-setosa
         125    Iris-virginica
         7      Iris-setosa
Name: class, dtype: object
```

```
In [302... # Comparing the mapped label names to the actual labels
[y_test_dict[i] for i in model_kmeans.predict(X_test2)] == y_test2
```

```
Out[302]: 17      True
          4       True
          61      True
          107     True
          20      True
          123     False
          43      True
          65      True
          93      True
          106     False
          137     True
          116     True
          47      True
          57      True
          42      True
          118     True
          132     True
          97      True
          50      True
          90      True
          92      True
          48      True
          104     True
          45      True
          125     True
          7       True
Name: class, dtype: bool
```